

Assignment 5.1

April 18, 2021

0.1 Assignment 5.1

Implement the movie review classifier found in section 3.4 of Deep Learning with Python.

```
[1]: # Import libraries
from tensorflow import keras
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop
```

```
[2]: import keras
keras.__version__
```

```
[2]: '2.4.3'
```

```
[4]: # Load the dataset
from keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.
    ↳load_data(num_words=10000)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>

17465344/17464789 [=====] - 1s 0us/step

<__array_function__ internals>:5: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

/opt/conda/lib/python3.8/site-packages/tensorflow/python/keras/datasets/imdb.py:159:

VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

x_train, y_train = np.array(xs[:idx]), np.array(labels[:idx])

/opt/conda/lib/python3.8/site-packages/tensorflow/python/keras/datasets/imdb.py:160:

VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences

(which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray

```
x_test, y_test = np.array(xs[idx:]), np.array(labels[idx:])
```

The argument num_words=10000 means that we will only keep the top 10,000 most frequently occurring words in the training data. Rare words will be discarded. This allows us to work with vector data of manageable size. The variables train_data and test_data are lists of reviews, each review being a list of word indices (encoding a sequence of words). train_labels and test_labels are lists of 0s and 1s, where 0 stands for “negative” and 1 stands for “positive”

```
[5]: train_data[0]
```

```
[5]: [1,
      14,
      22,
      16,
      43,
      530,
      973,
      1622,
      1385,
      65,
      458,
      4468,
      66,
      3941,
      4,
      173,
      36,
      256,
      5,
      25,
      100,
      43,
      838,
      112,
      50,
      670,
      2,
      9,
      35,
      480,
      284,
      5,
      150,
      4,
      172,
```

112,
167,
2,
336,
385,
39,
4,
172,
4536,
1111,
17,
546,
38,
13,
447,
4,
192,
50,
16,
6,
147,
2025,
19,
14,
22,
4,
1920,
4613,
469,
4,
22,
71,
87,
12,
16,
43,
530,
38,
76,
15,
13,
1247,
4,
22,
17,
515,
17,

12,
16,
626,
18,
2,
5,
62,
386,
12,
8,
316,
8,
106,
5,
4,
2223,
5244,
16,
480,
66,
3785,
33,
4,
130,
12,
16,
38,
619,
5,
25,
124,
51,
36,
135,
48,
25,
1415,
33,
6,
22,
12,
215,
28,
77,
52,
5,
14,

407,
16,
82,
2,
8,
4,
107,
117,
5952,
15,
256,
4,
2,
7,
3766,
5,
723,
36,
71,
43,
530,
476,
26,
400,
317,
46,
7,
4,
2,
1029,
13,
104,
88,
4,
381,
15,
297,
98,
32,
2071,
56,
26,
141,
6,
194,
7486,
18,

```
4,  
226,  
22,  
21,  
134,  
476,  
26,  
480,  
5,  
144,  
30,  
5535,  
18,  
51,  
36,  
28,  
224,  
92,  
25,  
104,  
4,  
226,  
65,  
16,  
38,  
1334,  
88,  
12,  
16,  
283,  
5,  
16,  
4472,  
113,  
103,  
32,  
15,  
16,  
5345,  
19,  
178,  
32]
```

```
[6]: train_labels[0]
```

```
[6]: 1
```

```
[7]: #Since we restricted ourselves to the top 10,000 most frequent words, no word_
      →index will exceed 10,000:
      max([max(sequence) for sequence in train_data])
```

[7]: 9999

Lastly, we need to pick a loss function and an optimizer. Since we are facing a binary classification problem and the output of our network is a probability (we end our network with a single-unit layer with a sigmoid activation), is it best to use the `binary_crossentropy` loss. It isn't the only viable choice: you could use, for instance, `mean_squared_error`. But crossentropy is usually the best choice when you are dealing with models that output probabilities. Crossentropy is a quantity from the field of Information Theory, that measures the "distance" between probability distributions, or in our case, between the ground-truth distribution and our predictions. Here's the step where we configure our model with the `rmsprop` optimizer and the `binary_crossentropy` loss function. Note that we will also monitor accuracy during training.

```
[8]: # word_index is a dictionary mapping words to an integer index
word_index = imdb.get_word_index()
# We reverse it, mapping integer indices to words
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])
# We decode the review; note that our indices were offset by 3
# because 0, 1 and 2 are reserved indices for "padding", "start of sequence",
→and "unknown".
decoded_review = ' '.join([reverse_word_index.get(i - 3, '?') for i in
      →train_data[0]]])
decoded_review
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb_word_index.json
1646592/1641221 [=====] - 0s 0us/step

```
[8]: "? this film was just brilliant casting location scenery story direction
     everyone's really suited the part they played and you could just imagine being
     there robert ? is an amazing actor and now the same being director ? father came
     from the same scottish island as myself so i loved the fact there was a real
     connection with this film the witty remarks throughout the film were great it
     was just brilliant so much that i bought the film as soon as it was released for
     ? and would recommend it to everyone to watch and the fly fishing was amazing
     really cried at the end it was so sad and you know what they say if you cry at a
     film it must have been good and this definitely was also ? to the two little
     boy's that played the ? of norman and paul they were just brilliant children are
     often left out of the ? list i think because the stars that play them all grown
     up are such a big profile for the whole film but these children are amazing and
     should be praised for what they have done don't you think the whole story was so
     lovely because it was true and was someone's life after all that was shared with
     us all"
```

Preparing the data: We cannot feed lists of integers into a neural network. We have to turn our

lists into tensors. There are two ways we could do that: We could pad our lists so that they all have the same length, and turn them into an integer tensor of shape (samples, word_indices), then use as first layer in our network a layer capable of handling such integer tensors (the Embedding layer, which we will cover in detail later in the book). We could one-hot-encode our lists to turn them into vectors of 0s and 1s. Concretely, this would mean for instance turning the sequence [3, 5] into a 10,000-dimensional vector that would be all-zeros except for indices 3 and 5, which would be ones. Then we could use as first layer in our network a Dense layer, capable of handling floatingpoint vector data. We will go with the latter solution. Let's vectorize our data, which we will do manually for maximum clarity:

```
[10]: import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    # Create an all-zero matrix of shape (len(sequences), dimension)
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1. # set specific indices of results[i] to 1s
    return results
# Our vectorized training data
x_train = vectorize_sequences(train_data)
# Our vectorized test data
x_test = vectorize_sequences(test_data)
```

```
[11]: # Print sample
x_train[0]
```

```
[11]: array([0., 1., 1., ..., 0., 0., 0.])
```

```
[12]: # Our vectorized labels
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')
```

Now our data is ready to be fed into a neural network. ##### Building our network #####
The input data is simply vectors, and the labels are scalars (1s and 0s). This is the easiest setup you will ever encounter. A type of network that performs well on such a problem would be a simple stack of fully connected (Dense) layers with relu activations:

Dense(16, activation='relu')

The argument being passed to each Dense layer (16) is the number of “hidden units” of the layer. What's a hidden unit? It's a dimension in the representation space of the layer. Each such Dense layer with a relu activation implements the following chain of tensor operations:

output = relu(dot(W, input) + b)

Having 16 hidden units means that the weight matrix W will have shape (input_dimension, 16), i.e. the dot product with W will project the input data onto a 16-dimensional representation space (and then we would add the bias vector b and apply the relu operation). You can intuitively understand the dimensionality of your representation space as “how much freedom you are allowing the network to have when learning internal representations”. Having more hidden units (a higher dimensional representation space) allows your network to learn more complex representations, but it makes your network more computationally expensive and may lead to learning unwanted patterns (patterns that will improve performance on the training data but not on the test data). There are

two key architecture decisions to be made about such stack of dense layers: How many layers to use. How many “hidden units” to choose for each layer.

```
[13]: #The Keras implementation
from keras import models
from keras import layers
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

Lastly, we need to pick a loss function and an optimizer. Since we are facing a binary classification problem and the output of our network is a probability (we end our network with a single-unit layer with a sigmoid activation), is it best to use the `binary_crossentropy` loss. It isn't the only viable choice: you could use, for instance, `mean_squared_error`. But crossentropy is usually the best choice when you are dealing with models that output probabilities. Crossentropy is a quantity from the field of Information Theory, that measures the “distance” between probability distributions, or in our case, between the ground-truth distribution and our predictions. Here's the step where we configure our model with the `rmsprop` optimizer and the `binary_crossentropy` loss function. Note that we will also monitor accuracy during training.

```
[14]: model.compile(optimizer='rmsprop',
    loss='binary_crossentropy',
    metrics=['accuracy'])
```

We are passing our optimizer, loss function and metrics as strings, which is possible because `rmsprop`, `binary_crossentropy` and `accuracy` are packaged as part of Keras. Sometimes you may want to configure the parameters of your optimizer, or pass a custom loss function or metric function. This former can be done by passing an optimizer class instance as the optimizer argument:

```
[15]: from keras import optimizers
model.compile(optimizer=optimizers.RMSprop(lr=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy'])
from keras import losses
from keras import metrics
model.compile(optimizer=optimizers.RMSprop(lr=0.001),
    loss=losses.binary_crossentropy,
    metrics=[metrics.binary_accuracy])
```

Validating our approach: In order to monitor during training the accuracy of the model on data that it has never seen before, we will create a “validation set” by setting apart 10,000 samples from the original training data:

```
[16]: x_val = x_train[:10000]
    partial_x_train = x_train[10000:]
    y_val = y_train[:10000]
    partial_y_train = y_train[10000:]
```

We will now train our model for 20 epochs (20 iterations over all samples in the `x_train` and `y_train` tensors), in mini-batches of 512 samples. At this same time we will monitor loss and accuracy on the 10,000 samples that we set apart. This is done by passing the validation data as the `validation_data` argument:

```
[17]: history = model.fit(partial_x_train,
    partial_y_train,
    epochs=20,
    batch_size=512,
    validation_data=(x_val, y_val))
```

Epoch 1/20

30/30 [=====] - 2s 45ms/step - loss: 0.6064 -
binary_accuracy: 0.7036 - val_loss: 0.4138 - val_binary_accuracy: 0.8627

Epoch 2/20

30/30 [=====] - 0s 13ms/step - loss: 0.3494 -
binary_accuracy: 0.8988 - val_loss: 0.3342 - val_binary_accuracy: 0.8723

Epoch 3/20

30/30 [=====] - 0s 12ms/step - loss: 0.2492 -
binary_accuracy: 0.9207 - val_loss: 0.2857 - val_binary_accuracy: 0.8869

Epoch 4/20

30/30 [=====] - 0s 12ms/step - loss: 0.1920 -
binary_accuracy: 0.9391 - val_loss: 0.2724 - val_binary_accuracy: 0.8899

Epoch 5/20

30/30 [=====] - 0s 13ms/step - loss: 0.1494 -
binary_accuracy: 0.9546 - val_loss: 0.2919 - val_binary_accuracy: 0.8833

Epoch 6/20

30/30 [=====] - 0s 12ms/step - loss: 0.1220 -
binary_accuracy: 0.9640 - val_loss: 0.2923 - val_binary_accuracy: 0.8857

Epoch 7/20

30/30 [=====] - 0s 13ms/step - loss: 0.1006 -
binary_accuracy: 0.9712 - val_loss: 0.3034 - val_binary_accuracy: 0.8859

Epoch 8/20

30/30 [=====] - 0s 12ms/step - loss: 0.0821 -
binary_accuracy: 0.9778 - val_loss: 0.3399 - val_binary_accuracy: 0.8793

Epoch 9/20

30/30 [=====] - 0s 12ms/step - loss: 0.0709 -
binary_accuracy: 0.9801 - val_loss: 0.3583 - val_binary_accuracy: 0.8760

Epoch 10/20

30/30 [=====] - 0s 12ms/step - loss: 0.0603 -
binary_accuracy: 0.9845 - val_loss: 0.3718 - val_binary_accuracy: 0.8794

Epoch 11/20

30/30 [=====] - 0s 12ms/step - loss: 0.0465 -
binary_accuracy: 0.9901 - val_loss: 0.3917 - val_binary_accuracy: 0.8766

Epoch 12/20

30/30 [=====] - 0s 12ms/step - loss: 0.0355 -
binary_accuracy: 0.9933 - val_loss: 0.4277 - val_binary_accuracy: 0.8753

Epoch 13/20

```

30/30 [=====] - 1s 19ms/step - loss: 0.0281 -
binary_accuracy: 0.9952 - val_loss: 0.4573 - val_binary_accuracy: 0.8726
Epoch 14/20
30/30 [=====] - 1s 38ms/step - loss: 0.0240 -
binary_accuracy: 0.9962 - val_loss: 0.4903 - val_binary_accuracy: 0.8723
Epoch 15/20
30/30 [=====] - 1s 34ms/step - loss: 0.0175 -
binary_accuracy: 0.9978 - val_loss: 0.5272 - val_binary_accuracy: 0.8685
Epoch 16/20
30/30 [=====] - 0s 17ms/step - loss: 0.0140 -
binary_accuracy: 0.9985 - val_loss: 0.5609 - val_binary_accuracy: 0.8692
Epoch 17/20
30/30 [=====] - 0s 13ms/step - loss: 0.0112 -
binary_accuracy: 0.9989 - val_loss: 0.6110 - val_binary_accuracy: 0.8644
Epoch 18/20
30/30 [=====] - 0s 16ms/step - loss: 0.0077 -
binary_accuracy: 0.9995 - val_loss: 0.6441 - val_binary_accuracy: 0.8669
Epoch 19/20
30/30 [=====] - 0s 15ms/step - loss: 0.0052 -
binary_accuracy: 0.9997 - val_loss: 0.7203 - val_binary_accuracy: 0.8612
Epoch 20/20
30/30 [=====] - 0s 11ms/step - loss: 0.0044 -
binary_accuracy: 0.9998 - val_loss: 0.7110 - val_binary_accuracy: 0.8649

```

```

[18]: history_dict = history.history
      history_dict.keys()
      #dict_keys(['loss', 'val_loss', 'binary_accuracy', 'val_binary_accuracy'])
      #history_dict = history.history
      print(history_dict)

```

```


{'loss': [0.5372720956802368, 0.32974353432655334, 0.23959974944591522,
0.1874733716249466, 0.15123368799686432, 0.12618596851825714,
0.10534216463565826, 0.08587293326854706, 0.0715474784374237,
0.06055949255824089, 0.049735452979803085, 0.041049495339393616,
0.03138168528676033, 0.025183064863085747, 0.020492268726229668,
0.014615150168538094, 0.012248418293893337, 0.009191435761749744,
0.008731354959309101, 0.003938845358788967], 'binary_accuracy':
[0.7860666513442993, 0.8981999754905701, 0.9225999712944031, 0.9380666613578796,
0.9517333507537842, 0.9595999717712402, 0.9675333499908447, 0.9750666618347168,
0.979533314704895, 0.9832000136375427, 0.98826664686203, 0.9905999898910522,
0.9933333396911621, 0.9958000183105469, 0.9961333274841309, 0.9983999729156494,
0.9984666705131531, 0.9991999864578247, 0.9983999729156494, 0.9998666644096375],
'val_loss': [0.41378292441368103, 0.3341832458972931, 0.28574100136756897,
0.2724001407623291, 0.29193294048309326, 0.29225224256515503,
0.3034060597419739, 0.3398527204990387, 0.358333557844162, 0.3718472421169281,
0.3917437493801117, 0.42767104506492615, 0.45734837651252747,
0.4902653396129608, 0.5272461771965027, 0.5609274506568909, 0.6109947562217712,
0.644079327583313, 0.7203235626220703, 0.7110231518745422],

```


```
'val_binary_accuracy': [0.8626999855041504, 0.8723000288009644,  
0.886900007724762, 0.8899000287055969, 0.8833000063896179, 0.885699987411499,  
0.8859000205993652, 0.8792999982833862, 0.8759999871253967, 0.8794000148773193,  
0.8766000270843506, 0.8752999901771545, 0.8726000189781189, 0.8723000288009644,  
0.8684999942779541, 0.8691999912261963, 0.8644000291824341, 0.8669000267982483,  
0.8611999750137329, 0.8648999929428101]]}
```

It contains 4 entries: one per metric that was being monitored, during training and during validation. Let's use Matplotlib to plot the training and validation loss side by side, as well as the training and validation accuracy:

```
[19]: import matplotlib.pyplot as plt  
acc = history.history['binary_accuracy']  
val_acc = history.history['val_binary_accuracy']  
#acc = history.history['acc']  
#val_acc = history.history['val_acc']  
loss = history.history['loss']  
val_loss = history.history['val_loss']  
epochs = range(1, len(acc) + 1)  
# "bo" is for "blue dot"  
plt.plot(epochs, loss, 'bo', label='Training loss')  
# b is for "solid blue line"  
plt.plot(epochs, val_loss, 'b', label='Validation loss')  
plt.title('Training and validation loss')  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.legend()  
plt.show()
```



```
[20]: plt.clf() # clear figure
      #acc_values = history_dict['acc']
      #val_acc_values = history_dict['val_acc']
      acc = history.history['binary_accuracy']
      val_acc = history.history['val_binary_accuracy']
      plt.plot(epochs, acc, 'bo', label='Training acc')
      plt.plot(epochs, val_acc, 'b', label='Validation acc')
      plt.title('Training and validation accuracy')
      plt.xlabel('Epochs')
      plt.ylabel('Loss')
      plt.legend()
      plt.show()
```



The dots are the training loss and accuracy, while the solid lines are the validation loss and accuracy. Note that your own results may vary slightly due to a different random initialization of your network. As you can see, the training loss decreases with every epoch and the training accuracy increases with every epoch. That’s what you would expect when running gradient descent optimization – the quantity you are trying to minimize should get lower with every iteration. But that isn’t the case for the validation loss and accuracy: they seem to peak at the fourth epoch. This is an example of what we were warning against earlier: a model that performs better on the training data isn’t necessarily a model that will do better on data it has never seen before. In precise terms, what you are seeing is “overfitting”: after the second epoch, we are over-optimizing on the training data, and we ended up learning representations that are specific to the training data and do not generalize to data outside of the training set. In this case, to prevent overfitting, we could simply stop training after three epochs. In general, there is a range of techniques you can leverage to mitigate overfitting. Let’s train a new network from scratch for four epochs, then evaluate it on our test data:

```
[21]: model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
loss='binary_crossentropy',
metrics=['accuracy'])
model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

```

Epoch 1/4
49/49 [=====] - 1s 10ms/step - loss: 0.5497 - accuracy:
0.7387
Epoch 2/4
49/49 [=====] - 0s 10ms/step - loss: 0.2711 - accuracy:
0.9051
Epoch 3/4
49/49 [=====] - 0s 8ms/step - loss: 0.2010 - accuracy:
0.9303
Epoch 4/4
49/49 [=====] - 0s 7ms/step - loss: 0.1674 - accuracy:
0.9435
782/782 [=====] - 2s 2ms/step - loss: 0.3180 -
accuracy: 0.8735

```

```
[22]: results
```

```
[22]: [0.3180374503135681, 0.8734800219535828]
```

Our fairly naive approach achieves an accuracy of 87%. With state-of-the-art approaches, one should be able to get close to 95%.

Using a trained network to generate predictions on new data After having trained a network, you will want to use it in a practical setting. You can generate the likelihood of reviews being positive by using the predict method:

```
[23]: model.predict(x_test)
```

```
[23]: array([[0.14494497],
            [0.999773  ],
            [0.56950504],
            ...,
            [0.07366014],
            [0.04316697],
            [0.29795432]], dtype=float32)
```

As you can see, the network is very confident for some samples (0.99 or more, or 0.01 or less) but less confident for others (0.6, 0.4).

Conclusion There's usually quite a bit of preprocessing you need to do on your raw data in order to be able to feed it – as tensors – into a neural network. In the case of sequences of words, they can be encoded as binary vectors – but there are other encoding options too. Stacks of Dense layers with relu activations can solve a wide range of problems (including sentiment classification), and you will likely use them frequently. In a binary classification problem (two output classes), your network should end with a Dense layer with 1 unit and a sigmoid activation, i.e. the output of your network should be a scalar between 0 and 1, encoding a probability. With such a scalar sigmoid output, on a binary classification problem, the loss function you should use is `binary_crossentropy`. The `rmsprop` optimizer is generally a good enough choice of optimizer, whatever your problem. That's one less thing for you to worry about.

Reference:

<https://github.com/fchollet/deep-learning-with-python-notebooks>

[]: